

# Surveillance Problem with Deep Learning

Emily Nguyen, Rui Fang

Oden Institute

2021

# Table of Contents

1 Surveillance Problem

2 Setup

3 Results

# Table of Contents

1 Surveillance Problem

2 Setup

3 Results

# Surveillance Problem


- With a known environment (comprised of free and occluded space), what is the minimum number of (and positions of the minima of) sensors to map out the environment?

# Surveillance Problem

- With a known environment (comprised of free and occluded space), what is the minimum number of (and positions of the minima of) sensors to map out the environment?
- Example scenarios
  - Placing lights in a city block to make every part visible
  - The minimum set of guards required to watch over an art gallery

# Surveillance Problem

- With a known environment (comprised of free and occluded space), what is the minimum number of (and positions of the minima of) sensors to map out the environment?
- Example scenarios
  - Placing lights in a city block to make every part visible
  - The minimum set of guards required to watch over an art gallery



# Problem Statement

- Consider a domain  $\Omega \subset \mathbb{R}^n$ . Let the domain be partitioned as  $\Omega = \Omega_{\text{free}} \cup \Omega_{\text{obs}}$ 
  - $\Omega_{\text{free}}$ : free space
  - $\Omega_{\text{obs}}$ : obstacles/occluded space with no holes

# Problem Statement

- Consider a domain  $\Omega \subset \mathbb{R}^n$ . Let the domain be partitioned as  $\Omega = \Omega_{\text{free}} \cup \Omega_{\text{obs}}$ 
  - $\Omega_{\text{free}}$ : free space
  - $\Omega_{\text{obs}}$ : obstacles/occluded space with no holes
- Find a sequence of vantage points  $x_i$  that maximize a gain function
  - $\mathcal{V}_{x_i}$ : visibility set (points visible from  $x_i$ )
  - $\Omega_k = \bigcup_{i=0}^k \mathcal{V}_{x_i}$ : cumulative visibility set (note:  $\Omega_{k-1} \subset \Omega_k$ )



# Problem Statement

- Consider a domain  $\Omega \subset \mathbb{R}^n$ . Let the domain be partitioned as  $\Omega = \Omega_{\text{free}} \cup \Omega_{\text{obs}}$ 
  - $\Omega_{\text{free}}$ : free space
  - $\Omega_{\text{obs}}$ : obstacles/occluded space with no holes
- Find a sequence of vantage points  $x_i$  that maximize a gain function
  - $\mathcal{V}_{x_i}$ : visibility set (points visible from  $x_i$ )
  - $\Omega_k = \bigcup_{i=0}^k \mathcal{V}_{x_i}$ : cumulative visibility set (note:  $\Omega_{k-1} \subset \Omega_k$ )
- The surveillance problem is: for a set of vantage points  $O$

$$\min_{O \subseteq \Omega_{\text{free}}} |O| \text{ subject to } \Omega_{\text{free}} = \bigcup_{x \in O} \mathcal{V}_x$$

# Problem Statement

- Consider a domain  $\Omega \subset \mathbb{R}^n$ . Let the domain be partitioned as  $\Omega = \Omega_{\text{free}} \cup \Omega_{\text{obs}}$ 
  - $\Omega_{\text{free}}$ : free space
  - $\Omega_{\text{obs}}$ : obstacles/occluded space with no holes
- Find a sequence of vantage points  $x_i$  that maximize a gain function
  - $\mathcal{V}_{x_i}$ : visibility set (points visible from  $x_i$ )
  - $\Omega_k = \bigcup_{i=0}^k \mathcal{V}_{x_i}$ : cumulative visibility set (note:  $\Omega_{k-1} \subset \Omega_k$ )
- The surveillance problem is: for a set of vantage points  $O$

$$\min_{O \subseteq \Omega_{\text{free}}} |O| \text{ subject to } \Omega_{\text{free}} = \bigcup_{x \in O} \mathcal{V}_x$$

- An implicit assumption: the sensor range is larger than the area of the environment

# Greedy Algorithm

- Can use a greedy algorithm based on choosing subsequent  $x_{k+1}$  that maximize information gain. Let the gain function be

$$g(x; \Omega_k) := |\mathcal{V}_x \cup \Omega_k| - |\Omega_k|$$

i.e., the volume of region visible from  $x$  but not from previous points  $x_0, x_1, \dots, x_k$

# Greedy Algorithm

- Can use a greedy algorithm based on choosing subsequent  $x_{k+1}$  that maximize information gain. Let the gain function be

$$g(x; \Omega_k) := |\mathcal{V}_x \cup \Omega_k| - |\Omega_k|$$

i.e., the volume of region visible from  $x$  but not from previous points  $x_0, x_1, \dots, x_k$

- We choose

$$x_{k+1} = \arg \max_{x \in \Omega_{\text{free}}} g(x; \Omega_k)$$

# Greedy Algorithm

- Can use a greedy algorithm based on choosing subsequent  $x_{k+1}$  that maximize information gain. Let the gain function be

$$g(x; \Omega_k) := |\mathcal{V}_x \cup \Omega_k| - |\Omega_k|$$

i.e., the volume of region visible from  $x$  but not from previous points  $x_0, x_1, \dots, x_k$

- We choose

$$x_{k+1} = \arg \max_{x \in \Omega_{\text{free}}} g(x; \Omega_k)$$

- However, this is very costly to compute

# Approximating Greedy Algorithm

- Try to train a model to predict the gains


# Approximating Greedy Algorithm

- Try to train a model to predict the gains
- Let the learned function be

$$g_{\theta}(x; \Omega_k, B_k)$$









where  $B_k$  is the part of  $\partial\Omega_k$  that might be in  $\Omega_{\text{free}}$  (i.e.,  $B_k = \partial\Omega_k \setminus \Omega_{\text{obs}}$ )

# Cumulative Visibility, Frontiers, and Gain Function





# Example Results



# Table of Contents

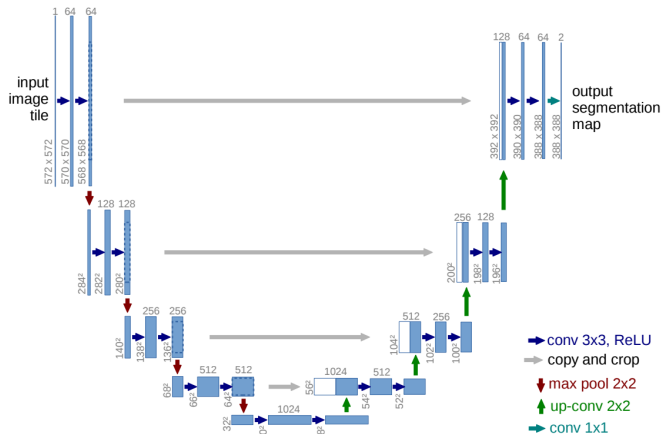
1 Surveillance Problem

2 Setup

3 Results

# U-Net

- Base the model on U-Net because we want the value predicted at each pixel to depend on a neighborhood



Source: <https://arxiv.org/pdf/1505.04597.pdf>

- Used two types of environments
  - Circle:  $64 \times 64$  pixel environments consisting of 2 to 6 circles
  - City:  $128 \times 128$  pixel aerial city blocks cropped from INRIA Aerial Image Labeling Dataset

- Used two types of environments
  - Circle:  $64 \times 64$  pixel environments consisting of 2 to 6 circles
  - City:  $128 \times 128$  pixel aerial city blocks cropped from INRIA Aerial Image Labeling Dataset
- Generated training data using the greedy algorithm


# Dataset

- Used two types of environments
  - Circle:  $64 \times 64$  pixel environments consisting of 2 to 6 circles
  - City:  $128 \times 128$  pixel aerial city blocks cropped from INRIA Aerial Image Labeling Dataset
- Generated training data using the greedy algorithm
- Removed data with very small gain, then normalized gain such that the maximum gain is 1 for each image


# Dataset

- Used two types of environments
  - Circle:  $64 \times 64$  pixel environments consisting of 2 to 6 circles
  - City:  $128 \times 128$  pixel aerial city blocks cropped from INRIA Aerial Image Labeling Dataset
- Generated training data using the greedy algorithm
- Removed data with very small gain, then normalized gain such that the maximum gain is 1 for each image
- Got  $\sim 50,000$  data pairs for Circle, and  $\sim 20,000$  data pairs for City


# Inputs and Targets (Circle)




(a) Cumulative Visibility 1




(b) Frontier 1




(c) Gain 1




(d) Cumulative Visibility 2




(e) Frontier 2




(f) Gain 2



(g) Cumulative Visibility 3




(h) Frontier 3



(i) Gain 3



# Inputs and Targets (City)



# Training

- Trained a model for each of Circle and City, used 90% of data for training, 10% for validation

# Training

- Trained a model for each of Circle and City, used 90% of data for training, 10% for validation
- Loss function is binary cross entropy

$$L(p, q) = -[p \log(q) + (1 - p) \log(1 - q)]$$

# Training

- Trained a model for each of Circle and City, used 90% of data for training, 10% for validation
- Loss function is binary cross entropy

$$L(p, q) = -[p \log(q) + (1 - p) \log(1 - q)]$$

- Optimization method is Adam
  - Initial learning rate:  $10^{-3}$

# Training

- Trained a model for each of Circle and City, used 90% of data for training, 10% for validation
- Loss function is binary cross entropy

$$L(p, q) = -[p \log(q) + (1 - p) \log(1 - q)]$$

- Optimization method is Adam
  - Initial learning rate:  $10^{-3}$
- Used PyTorch and Google Colab

# Table of Contents

1 Surveillance Problem

2 Setup

3 Results


# Results: Circular Environments

- These training outputs are from 1770 epochs in, which is indicated to be the point where the losses start to stagnate


# Results: Circular Environments

- These training outputs are from 1770 epochs in, which is indicated to be the point where the losses start to stagnate


expected training gain is:




output training gain is:




expected training gain is:




output training gain is:



expected training gain is:



output training gain is:






# Results: Circular Environments

- However, there are some environments that the model does not predict as well


# Results: Circular Environments

- However, there are some environments that the model does not predict as well


expected training gain is:




output training gain is:




expected training gain is:




output training gain is:





expected training gain is:



output training gain is:



# Results: Circular Environments Losses




# Results: Circular Environments Generalization

- The model seems to generalize well on the validation set


# Results: Circular Environments Generalization

- The model seems to generalize well on the validation set


expected test gain is:




output test gain is:




expected test gain is:




output test gain is:



expected test gain is:




output test gain is:




# Results: City Environments

- Loss vs Epoch for 500 epochs



# Results: City Environments

- Loss vs Epoch for 500 epochs




- Learning rate was decreased by a factor of 10 at epoch 50, 100, 150, resulting in slowed down decreases in losses
- The model doesn't generalize well on validation set


# Results: City Environments

- Prediction on validation set
  - Good examples

Normalization factor: 0.0031862073568201706



Normalization factor: 0.022148026748628015







# Results: City Environments

- Prediction on validation set
  - Bad examples

Normalization factor: 0.17189977251917601



Normalization factor: 0.005362153844404677



# Discussion

- For the circular environment, the model seems to combine circles if they seem too similar in shape to a bigger circle

# Discussion

- For the circular environment, the model seems to combine circles if they seem too similar in shape to a bigger circle
- A disadvantage of using Google Colab is the runtime limit where the runtime disconnects after a certain number of hours. This makes restarting more difficult in terms of keeping track of the learning rate
  - The validation and training losses experience an unusual jump upon continuing
  - A fix for this will be to train locally or on TACC

# Discussion

- For the circular environment, the model seems to combine circles if they seem too similar in shape to a bigger circle
- A disadvantage of using Google Colab is the runtime limit where the runtime disconnects after a certain number of hours. This makes restarting more difficult in terms of keeping track of the learning rate
  - The validation and training losses experience an unusual jump upon continuing
  - A fix for this will be to train locally or on TACC
- For the city environment, in many cases the map contains “isolated” free space, e.g., a courtyard surrounded by buildings. It is not going to be seen by any vantage point placed outside. Then it is extremely difficult for the network to predict the gain over such space.

# Discussion

- For the circular environment, the model seems to combine circles if they seem too similar in shape to a bigger circle
- A disadvantage of using Google Colab is the runtime limit where the runtime disconnects after a certain number of hours. This makes restarting more difficult in terms of keeping track of the learning rate
  - The validation and training losses experience an unusual jump upon continuing
  - A fix for this will be to train locally or on TACC
- For the city environment, in many cases the map contains “isolated” free space, e.g., a courtyard surrounded by buildings. It is not going to be seen by any vantage point placed outside. Then it is extremely difficult for the network to predict the gain over such space.
  - Can improve this by giving the network more direct information of the map, e.g., using cumulative visibility and obstacles as input, rather than using cumulative visibility and frontier

# Future Directions

- Tune hyperparameters of model and training procedure, in order to make the model more generalizable

# Future Directions

- Tune hyperparameters of model and training procedure, in order to make the model more generalizable
- Generate sequences of vantage points based on approximated gain function, and compare against the points generated with exact gain function

# Future Directions

- Tune hyperparameters of model and training procedure, in order to make the model more generalizable
- Generate sequences of vantage points based on approximated gain function, and compare against the points generated with exact gain function
- Frequency maps: for a certain map, using many sequences of vantage points generated with different initial conditions, see hot spots for where the greedy algorithm tends to put sensors/pick vantage points



# Future Directions

- Tune hyperparameters of model and training procedure, in order to make the model more generalizable
- Generate sequences of vantage points based on approximated gain function, and compare against the points generated with exact gain function
- Frequency maps: for a certain map, using many sequences of vantage points generated with different initial conditions, see hot spots for where the greedy algorithm tends to put sensors/pick vantage points
- Build a model to learn the “true” (unnormalized) gain function

# Future Directions

- Tune hyperparameters of model and training procedure, in order to make the model more generalizable
- Generate sequences of vantage points based on approximated gain function, and compare against the points generated with exact gain function
- Frequency maps: for a certain map, using many sequences of vantage points generated with different initial conditions, see hot spots for where the greedy algorithm tends to put sensors/pick vantage points
- Build a model to learn the “true” (unnormalized) gain function
  - Option 1: train the “true” gain function with MSE loss, with ReLU as the final layer
  - Option 2: create an additional module, MLP for example, connecting the bottom layer of the current Unet to a single output for the normalization factor

- Olaf Ronneberger and Philipp Fischer and Thomas Brox, U-Net: Convolutional Networks for Biomedical Image Segmentation
- Louis Ly thesis